

# linear regreesion

March 20, 2024

```
[8]: import numpy as np
import pandas as pd
import warnings
warnings.filterwarnings('ignore')
```

```
[9]: #importing data
data=pd.read_csv("C:\\Users\\user\\Downloads\\Dataset Heart Disease.csv")
data
```

```
[9]:
```

	age	chest pain type	resting bps	cholesterol	max heart rate
0	40	2	140	289	172
1	49	3	160	180	156
2	37	2	130	283	98
3	48	4	138	214	108
4	54	3	150	195	122
..	...	...	...	...	...
893	42	3	130	180	150
894	45	2	128	308	170
895	57	1	165	289	124
896	64	3	125	309	131
897	41	3	112	250	179

[898 rows x 5 columns]

calling arrays

```
[10]: x=np.array(data["age"]).reshape(-1,1)
```

```
[11]: y=np.array(data["max heart rate"])
```

checking for missing data


```
[12]: data.isna().sum()
```

```
[12]: age                0
chest pain type        0
resting bps            0
cholesterol            0
max heart rate         0
```

dtype: int64

graph

```
[13]: import matplotlib.pyplot as plt
plt.scatter(x,y)
plt.xlabel("age")
plt.ylabel("max heart rate")
plt.title("Relationship between age and heart disease")
plt.show()
```



splitting data

```
[14]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.
↪2,random_state=42)
```

```
#standardizing data from sklearn.preprocessing import StandardScaler scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train) x_test_scaled = scaler.transform(x_test)
```

```
[15]: #building the model  
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model
```

```
[15]: LinearRegression()
```

```
[16]: #model fitting  
model.fit(x_train,y_train)
```

```
[16]: LinearRegression()
```

```
[17]: #making predictions  
y_pred=model.predict(x_test)
```

```
[18]: model.coef_
```

```
[18]: array([-0.96626539])
```

```
[19]: #finding intercept(regression line intercepts with the y axis)  
model.intercept_
```

```
[19]: 193.4557211223224
```

```
[20]: #model evaluation  
#model accuracy on train values  
model.score(x_train,y_train)
```

```
[20]: 0.1411214865128494
```

```
[21]: #model accuracy on the test values  
model.score(x_test,y_test)
```

```
[21]: 0.09956280781317806
```

```
[22]: from sklearn.metrics import mean_squared_error,r2_score,mean_absolute_error  
print(mean_squared_error(y_test,y_pred))
```

```
573.7417929335393
```

```
[23]: print(mean_absolute_error(y_test,y_pred))
```

```
19.89992588788073
```

```
[24]: print(r2_score(y_test,y_pred))
```

```
0.09956280781317806
```

MODEL OPTIMIZATION

```
[25]: from sklearn.model_selection import GridSearchCV
      model = LinearRegression()
```

```
[26]: #performing GridSearchCV to find best hyperparameters
      param_grid = {'fit_intercept': [True, False], 'positive': [True, False], 'positive':
        ↪ [True, False]}
      param_grid
```

```
[26]: {'fit_intercept': [True, False], 'positive': [True, False]}
```

```
[27]: #fitting data
      data= GridSearchCV(model,param_grid, cv=5)
      data.fit(x_train,y_train)
```

```
[27]: GridSearchCV(cv=5, estimator=LinearRegression(),
                  param_grid={'fit_intercept': [True, False],
                              'positive': [True, False]})
```

```
[28]: #getting the best parameter from the GridSearch
      best_params = data.best_params_['fit_intercept']
      best_params
```

```
[28]: True
```

```
[29]: best_params = data.best_params_['positive']
      best_params
```

```
[29]: False
```

```
[30]: #training the model using the best parameters
      model=LinearRegression(fit_intercept=best_params)
      model.fit(x_train,y_train)
```

```
[30]: LinearRegression(fit_intercept=False)
```

```
[31]: #making predictions
      y_pred = model.predict(x_test)
```

```
[32]: #evaluating model performance
      print(mean_absolute_error(y_test,y_pred))
```

```
32.93329029906092
```

```
[33]: print(r2_score(y_test,y_pred))
```

```
-1.5683088109567058
```

```
[34]: #finding accuracy on train values
      model.score(x_train,y_train)
```

[34]: -1.813959890359337

```
[35]: #finding accuracy on test values  
      model.score(x_test,y_test)
```

[35]: -1.5683088109567058